

# CS-E3210 Machine Learning: Basic Principles

## Lecture 9: Clustering slides by Alex Jung, 2017

Department of Computer Science  
Aalto University, School of Science

Autumn (Period I) 2017

# Today's Motto

Your Friends Probably Look Like You

# Outline

- 1 Introduction
- 2 Clustering is NOT Classification
- 3 Hard Clustering
- 4 Soft Clustering
- 5 Summary

## What you should learn today...

- difference between supervised and unsupervised learning
- organize datasets into clusters
- difference between hard and soft clustering
- one algorithm for hard clustering
- one algorithm for soft clustering


# Applications of Clustering

- biology: **group** homologous **sequences** into **gene families**
- marketing: **partition** consumers into market segments
- sociology: find **communities** in social networks
- natural language processing: identify **topics** in a corpus
- computer vision: group pixels to **segments**
- climatology: find **weather regimes**
- ...

# Applications of Clustering: Market Segmentation



# Applications of Clustering: Social Network Analysis




# Applications of Clustering: Text Document Analysis





## Applications of Clustering: Image Segmentation



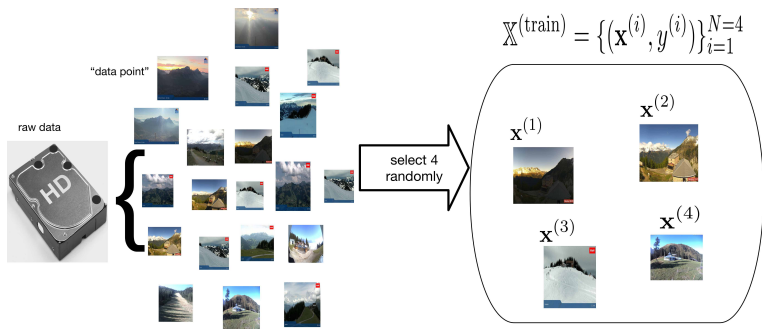
# Outline

- 1 Introduction
- 2 Clustering is NOT Classification
- 3 Hard Clustering
- 4 Soft Clustering
- 5 Summary

## Ski Resort Marketing

- you still did not find another job
- thus, you still work as marketing of a ski resort
- hard disk full of webcam snapshots (**gigabytes of data**)
- want to group them into “winter” and “summer” images
- you have only a **few hours** for this task ...

# The Dataset



## ML workflow so far...

- create  $\mathbb{X}^{(\text{train})} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^{N_t}$  by **manual labeling**
- **features**  $\mathbf{x}^{(i)} \in \mathcal{X}$  and **label**  $y^{(i)} \in \mathcal{Y}$  of  $i$ th data point
- define **loss**  $L((\mathbf{x}, y), h(\cdot))$  (e.g.,  $L((\mathbf{x}, y), h(\cdot)) = (y - h(\mathbf{x}))^2$ )
- define **hypothesis space**  $\mathcal{H}$  (e.g., linear maps  $h(\mathbf{x}) = \mathbf{w}^T \mathbf{x}$ )
- learn predictor  $h(\cdot) : \mathcal{X} \rightarrow \mathcal{Y}$  by **empirical risk minimization**

$$\min_{h(\cdot) \in \mathcal{H}} \mathcal{E}\{h(\cdot) | \mathbb{X}^{(\text{train})}\} = \min_{h(\cdot) \in \mathcal{H}} \frac{1}{|\mathbb{X}^{(\text{train})}|} \sum_{(\mathbf{x}, y) \in \mathbb{X}^{(\text{train})}} L((\mathbf{x}, y), h(\cdot))$$

## NO Time For Labeling

- already spent 3 weeks on grouping into winter/summer
- we have time for manual labelling only one picture
- can we cluster/group the snapshots directly into two groups ?
- if clustering works, need to look at ONE SINGLE snapshot

# How To Group or Cluster into Two Clusters?



## Definition of Clustering?

- **informal description** according to **Wikipedia**:

“Cluster analysis or clustering is the task of assigning a set of objects into groups (called clusters) so that the objects in the same cluster are more similar (in some sense or another) to each other than to those in other clusters.”

- **no single best formal definition** of clustering!



# Look At Data in Feature Space

# Clustering vs. Classification

- common: **feature vector**  $\mathbf{x} \in \mathbb{R}^d$  and **label**  $y \in \{0, 1\}$
- **classification** is **supervised learning** method
  - need **labeled training data**  $\mathbb{X}^{(\text{train})} = \{(\mathbf{x}^{(i)}, y^{(i)})\}_{i=1}^{N_{\text{train}}}$
  - learn classifier (LogReg, SVM,...) via ERM using  $\mathbb{X}^{(\text{train})}$
  - **predict** label  $y$  for (classify) **new** snapshot with features  $\mathbf{x}$
- **clustering** is **unsupervised learning** method
  - **unlabeled** data  $\mathbb{X} = \{\mathbf{x}^{(i)}\}_{i=1}^N$
  - **find cluster index**  $y^{(i)}$  for each data vector  $\mathbf{x}^{(i)}$
  - clustering based on **intrinsic geometry of data points**  $\mathbb{X}$

## Hard vs. Soft-Clustering

- hard clustering:
  - data points belong to **one and only one cluster**,  $y^{(i)} \in \{0, 1\}$
  - data points partitioned into non-overlapping clusters
  - hard-clustering method: **K-means**
- soft clustering:
  - datapoint may belong to **several clusters**
  - clusters are **overlapping**
  - **strength of association/degree of belonging**  $y^{(i)} \in [0, 1]$
  - soft-clustering using **Gaussian mixture models** (covered in APM course CS-E4820)

# Outline

- 1 Introduction
- 2 Clustering is NOT Classification
- 3 Hard Clustering**
- 4 Soft Clustering
- 5 Summary

# K-means Clustering

## The Basic Idea

- partition  $\mathbb{X} = \{\mathbf{x}^{(i)}\}_{i=1}^N$  into clusters  $\mathcal{C}_y$ ,  $y \in \{0, \dots, K-1\}$
- hard clustering: each  $\mathbf{x}^{(i)}$  belongs **exactly to one**  $\mathcal{C}_{y^{(i)}}$
- cluster  $\mathcal{C}_y$  represented by **cluster mean**  $\mathbf{m}_y$
- popular clustering method: **K-means**

# K-means Clustering

The Algorithm (also called Lloyd's Algorithm)

1. input:  $\mathbb{X} = \{\mathbf{x}^{(i)}\}_{i=1}^N$ , number  $K$ , cluster means  $\{\mathbf{m}_y\}_{y=0}^{K-1}$

2. repeat until convergence

- **cluster assignment**: for each  $\mathbf{x}^{(i)}$  find nearest cluster mean

$$y^{(i)} = \underset{y \in \{0, \dots, K-1\}}{\operatorname{argmin}} \|\mathbf{x}^{(i)} - \mathbf{m}_y\|_2$$

- **mean update**: for each cluster  $\mathcal{C}_y = \{\mathbf{x}^{(i)} : y^{(i)} = y\}$ , compute

$$\mathbf{m}_y = (1/|\mathcal{C}_y|) \sum_{i: y^{(i)}=y} \mathbf{x}^{(i)}$$

- output: **cluster means**  $\{\mathbf{m}_y\}_{y=0}^{K-1}$  and **assignments**  $\{y^{(i)}\}_{i=1}^N$

# K-means Clustering

## The Algorithm in Action




Figure 9.1 of Bishop (2006)

# K-means Clustering

## Initialization

- $K$ -means needs initial choice for cluster means  $\mathbf{m}_y$
- no optimal choice in general; some heuristics:
  - use  $K$  randomly selected data points
  - use  $K$  random perturbations of sample mean
  - divide range of principal component into  $K$  grid points



# K-means Clustering

## The Optimization Problem

- K-means can be interpreted as **optimization method**
- for  $\{\mathbf{m}_y\}_{y=0}^{K-1}$ ,  $\{y^{(i)}\}_{i=1}^N$ , define **cost/distortion** function:

$$\mathcal{E}(\{\mathbf{m}_y\}_{y=0}^{K-1}, \{y^{(i)}\}_{i=1}^N) = \sum_{i=1}^N \left\| \mathbf{x}^{(i)} - \mathbf{m}_{y^{(i)}} \right\|^2$$

- $\mathcal{E}(\{\mathbf{m}_y\}_{y=0}^{K-1}, \{y^{(i)}\}_{i=1}^N)$  is **non-convex and non-smooth!**
- K-means=**coordinate descent** for  $\mathcal{E}(\{\mathbf{m}_y\}_{y=0}^{K-1}, \{y^{(i)}\}_{i=1}^N)$
- allows for **convergence diagnosis**

## Coordinate Descent

- consider function  $f(x, y)$  of two variables  $x, y$
- we aim for  $x_0, y_0$  such that  $f(x_0, y_0) = \min_{x, y} f(x, y)$
- this minimization problem is often difficult
- sometimes, the minimization of  $f(x, y)$  either w.r.t. to  $x$  and fixed  $y$  and vice-versa is easy
- coordinate descent:
  - given current guess  $x_k, y_k$  obtain new  $x_{k+1}$  by
$$x_{k+1} = \underset{x}{\operatorname{argmin}} f(x, y_k)$$
  - obtain new  $y_{k+1}$  by
$$y_k = \underset{y}{\operatorname{argmin}} f(x_{k+1}, y)$$

# K-means Clustering


## Convergence

- $K$ -means is coord. descent for  $\mathcal{E}(\{\mathbf{m}_y\}_{y=0}^{K-1}, \{y^{(i)}\}_{i=1}^N)$
- objective  $\mathcal{E}$  **monotonically decreasing** throughout iterations
- however,  $K$ -means can get stuck in **local minimum** of  $\mathcal{E}$
- workaround: **run  $K$ -means several times with random initial.**
- pick solution **yielding smallest cost  $\mathcal{E}$**

# K-means Clustering

## How To Choose $K$

- e.g., by finding “elbow” in distortion curve



- often clustering used as pre-processing for learning method
- choose  $K$  by cross-validation of overall method
- use complexity penalization (favouring smaller  $K$ )

# K-means Clustering

## Properties

- conceptually and algorithmically simple
- typically only small number of iterations required
- K-means **sensitive to initialization**
- iterations might get stuck in **local optimum**
- workaround: run K-means **several times with random init.**
- select solution with **smallest cost  $\mathcal{E}$**

# Outline

- 1 Introduction
- 2 Clustering is NOT Classification
- 3 Hard Clustering
- 4 Soft Clustering
- 5 Summary

## “Lets Put on the Probabilistic Glasses”

- lets consider  $K = 2$  for simplicity (extended easily to other  $K$ )
- lets interpret  $y^{(i)}$  as probability of  $\mathbf{x}^{(i)} \in \mathcal{C}_1$
- $y^{(i)} = P(\mathbf{x}^{(i)} \in \mathcal{C}_1 | \mathbb{X})$  “degree of  $\mathbf{x}^{(i)}$  belonging to  $\mathcal{C}_1$ ”
- what is degree of  $\mathbf{x}^{(i)}$  belonging to  $\mathcal{C}_0$ ?
- $K$ -means enforces  $P(\mathbf{x}^{(i)} \in \mathcal{C}_1 | \mathbb{X}) \in \{0, 1\}$





## Gaussian Mixture Model (GMM)

- cluster  $\mathcal{C}_y$  represented by Gaussian distribution  $\mathcal{N}(\mathbf{x}; \mathbf{m}_y, \mathbf{C}_y)$ <sup>1</sup>
- cluster  $\mathcal{C}_0$  has mean  $\mathbf{m}_0 \in \mathbb{R}^d$  and covariance  $\mathbf{C}_0 \in \mathbb{R}^{d \times d}$
- cluster  $\mathcal{C}_1$  has mean  $\mathbf{m}_1 \in \mathbb{R}^d$  and covariance  $\mathbf{C}_1 \in \mathbb{R}^{d \times d}$
- probability of data point  $\mathbf{x}^{(i)}$  belonging to  $\mathcal{C}_1$  is

$$y^{(i)} = \frac{\mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_1, \mathbf{C}_1)}{\mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_0, \mathbf{C}_0) + \mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_1, \mathbf{C}_1)} \in [0, 1] \quad \text{💬}$$

---

<sup>1</sup> $\mathcal{N}(\mathbf{x}; \mathbf{m}, \mathbf{C}) = \frac{\exp\left(-\frac{1}{2}(\mathbf{x}-\mathbf{m})^T \mathbf{C}^{-1}(\mathbf{x}-\mathbf{m})\right)}{\sqrt{\det\{2\pi\mathbf{C}\}}}$

## (Approximate) Maximum Likelihood

- consider current guess for  $y^{(i)}$  (degree of  $\mathbf{x}^{(i)} \in \mathcal{C}_1$ )
- “effective size” of  $\mathcal{C}_1$  is  $N_1 = \sum_{i=1}^N y^{(i)}$
- “effective size” of  $\mathcal{C}_0$  is  $N_0 = \sum_{i=1}^N (1 - y^{(i)}) = N - N_1$
- approximate  $\mathbf{m}_1$  by  $(1/N_1) \sum_{i=1}^N y^{(i)} \mathbf{x}^{(i)}$
- approx.  $\mathbf{C}_1$  by  $(1/N_1) \sum_{i=1}^N y^{(i)} (\mathbf{x}^{(i)} - \mathbf{m}_1) (\mathbf{x}^{(i)} - \mathbf{m}_1)^T$
- similarly for  $\mathbf{m}_0$  and  $\mathbf{C}_0$

## A Soft-Clustering Algorithm

- 1: use initial guess for GMM parameters  $\mathbf{m}_0, \mathbf{m}_1, \mathbf{C}_0, \mathbf{C}_1$

- 2: update degrees of belonging

$$y^{(i)} = \mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_1, \mathbf{C}_1) / (\mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_0, \mathbf{C}_0) + \mathcal{N}(\mathbf{x}^{(i)}; \mathbf{m}_1, \mathbf{C}_1))$$

- 3: update GMM parameters  $N_1 = \sum_{i=1}^N y^{(i)}, N_0 = N - N_1,$

$$\mathbf{m}_1 = \frac{1}{N_1} \sum_{i=1}^N y^{(i)} \mathbf{x}^{(i)}, \mathbf{C}_1 = \frac{1}{N_1} \sum_{i=1}^N y^{(i)} (\mathbf{x}^{(i)} - \mathbf{m}_1) (\mathbf{x}^{(i)} - \mathbf{m}_1)^T$$

$$\mathbf{m}_0 = \frac{1}{N_0} \sum_{i=1}^N (1 - y^{(i)}) \mathbf{x}^{(i)}, \mathbf{C}_0 = \frac{1}{N_0} \sum_{i=1}^N (1 - y^{(i)}) (\mathbf{x}^{(i)} - \mathbf{m}_0) (\mathbf{x}^{(i)} - \mathbf{m}_0)^T$$

- 4: if not converged go to step 2

# Soft Clustering




Figure 9.8 of Bishop (2006)

# A Soft Clustering Algorithm

## Properties

- based on **generative GMM** for data
- implicitly estimates GMM parameters ( $\mathbf{m}_0, \mathbf{C}_0, \dots$ )
- problem of **local optima** (use **several random initializations**)
- **soft cluster assignments** (**degree of belonging**)  $y^{(i)} \in [0, 1]$
- reduces to **K-means** for  $\mathbf{C}_0 = \mathbf{C}_1 = \sigma^2 \mathbf{I}$  with small  $\sigma^2$

# Outline

- 1 Introduction
- 2 Clustering is NOT Classification
- 3 Hard Clustering
- 4 Soft Clustering
- 5 Summary

# Summary

what we learned today ...

- difference between **soft-** and **hard clustering**
- one hard-clustering algorithm, i.e., ***K*-means**
- one **soft clustering** algorithm (Gaussian mixture models)

# What happens next?

- next lecture **feature learning**
- recommended preparation: read Chap 5.8. [DLBook]